
1. ENI – Elastic Network Interface

What It Is:

An **ENI (Elastic Network Interface)** is the basic network interface in AWS, functioning like a virtual network card attached to an EC2 instance. It includes:

- Private IP address
- Public IP address (if assigned)
- MAC address
- Security groups
- One or more secondary IPs

Use Cases:

- **Multi-homed instances:** Attach multiple ENIs for different subnets or security zones.
- **Network appliances:** Use ENIs for firewalls, routers, or proxies.
- **High availability failover:** Detach an ENI from a failed instance and reattach it to another.

Common Scenario:

A web server instance with a primary ENI for app traffic and a secondary ENI for admin traffic in a separate subnet.

2. ENA – Elastic Network Adapter

What It Is:

An **ENA (Elastic Network Adapter)** is a high-performance network interface used to provide **enhanced networking** for EC2 instances. It supports:

- Up to **100 Gbps** throughput
- Low latency
- Low jitter
- Single root I/O virtualization (**SR-IOV**)

Use Cases:

- **High-performance workloads:** Big data, video rendering, gaming servers
- **Instances in Cluster Placement Groups**
- **Burstable applications** that need predictable performance

Common Scenario:

A Spark cluster running on C5n instances uses ENA to achieve high-speed communication between nodes in a Cluster Placement Group.

Supported Instance Types:

- C5, M5, R5, and newer families
- Most Nitro-based EC2 instances

3. EFA – Elastic Fabric Adapter

What It Is:

An **EFA (Elastic Fabric Adapter)** is a specialized network interface for **HPC (High Performance Computing)** workloads. It extends ENA functionality to provide:

- Low latency
- High throughput
- **OS-bypass networking** (using libfabric API)
- Integrated with **MPI libraries** for parallel computing

Use Cases:

- **MPI-based HPC applications:** Computational fluid dynamics, genomics, seismic analysis
- **AI/ML training:** Especially on large distributed training jobs
- **Tightly coupled applications** that require extremely low latency

Common Scenario:

A machine learning training cluster using P4d instances with EFA interfaces to accelerate model training through ultra-low-latency node communication.

Notes:

- EFA must be **enabled at launch**.
- Supported only on **specific EC2 instance types** (e.g., C5n, P4d, M5n).
- Works **only within a placement group** and **same VPC/subnet**.

Summary Comparison Table:

Feature	ENI	ENA	EFA
Purpose	Basic networking	Enhanced networking	HPC and tightly coupled workloads
Performance	Standard	High throughput, low latency	Ultra-low latency, OS-bypass
Max Throughput	Up to 10 Gbps	Up to 100 Gbps	Up to 100 Gbps with low latency
Use Case	Web apps, NAT, proxies	Big data, web-scale apps, Cluster PG	HPC, AI/ML training, MPI apps
Instance Types	All	Nitro-based families (C5, M5, R5, etc.)	Select types only (e.g., P4d, C5n, M5n)
Special APIs	No	No	Yes (libfabric, MPI integration)
